

Haskell Beats C Using Generalized Stream Fusion

Geoffrey Mainland

Microsoft Research Ltd

Chalmers, December 2012

Joint work with Roman Leshchinskiy, Simon Peyton Jones, and Simon Marlow.

The challenge...

- ▶ Write high-level, *high-performance*, sequence-processing code. Think numerical code and vectors.
- ▶ In particular, allow compositional code without sacrificing performance.
- ▶ Should be able to take advantage of special-case operations, e.g., bulk memory copies and SIMD instructions.

The plan

- ▶ Overview of streams and stream fusion (old news).
- ▶ *Generalized* stream fusion. Simultaneously maintain several representations, and pick the best one depending on operations performed.
- ▶ Craft a representation tailored for computing with SSE instructions.
- ▶ Produce efficient numerical code from Haskell, i.e., how to make Haskell beat C!

Fusion

$\text{map } f \circ \text{map } g$

Fusion

$$\text{map } f \circ \text{map } g \equiv \text{map } (f \circ g)$$

Stream fusion

- ▶ Recursive functions are hard for the compiler to optimize.
- ▶ So turn recursive functions into non-recursive functions via a change in representation.
- ▶ Instead of working with lists or vectors, work with Streams.

Stream fusion

data Stream a **where**

Stream :: (s → Step s a) → s → Int → Stream a

data Step s a = Yield a s

| Skip s

| Done

map, now with Streams

$\text{map} :: (a \rightarrow b) \rightarrow \text{Vector } a \rightarrow \text{Vector } b$
 $\text{map } f = \text{unstream} \circ \text{map}_s f \circ \text{stream}$

map, now with Streams

$\text{map} :: (a \rightarrow b) \rightarrow \text{Vector } a \rightarrow \text{Vector } b$
 $\text{map } f = \text{unstream} \circ \text{map}_s f \circ \text{stream}$

$\text{map}_s :: (a \rightarrow b) \rightarrow \text{Stream } a \rightarrow \text{Stream } b$
 $\text{map}_s f (\text{Stream step } s) = \text{Stream step}' s$
where

$\text{step}' s = \text{case step } s \text{ of}$
 $\text{Yield } x \ s' \rightarrow \text{Yield } (f\ x) \ s'$
 $\text{Skip } s' \rightarrow \text{Skip } s'$
 $\text{Done} \rightarrow z$

Fusing maps

$$\begin{aligned} \text{map } f \circ \text{map } g &\equiv \\ \text{unstream} \circ \text{map}_s f \circ \text{stream} \circ \text{unstream} \circ \text{map}_s g \circ \text{stream} \end{aligned}$$

Fusing maps

$$\begin{aligned} \text{map } f \circ \text{map } g &\equiv \\ \text{unstream} \circ \text{map}_s f \circ \text{stream} \circ \text{unstream} \circ \text{map}_s g \circ \text{stream} \end{aligned}$$

$$\begin{aligned} \text{map } f \circ \text{map } g &\equiv \\ \text{unstream} \circ \text{map}_s f \circ \text{map}_s g \circ \text{stream} \end{aligned}$$

Fusing maps

$$\begin{aligned} \text{map } f \circ \text{map } g &\equiv \\ \text{unstream} \circ \text{map}_s f \circ \text{stream} \circ \text{unstream} \circ \text{map}_s g \circ \text{stream} \end{aligned}$$

$$\begin{aligned} \text{map } f \circ \text{map } g &\equiv \\ \text{unstream} \circ \text{map}_s f \circ \text{map}_s g \circ \text{stream} \end{aligned}$$

We get $\text{map } f \circ \text{map } g \equiv \text{map } (f \circ g)$ *for free!*

Fusing vector dot product

```
dotp :: Vector Double → Vector Double → Double  
dotp v w = sum (zipWith (*) v w)
```

Fusing vector dot product

$\text{dotp} :: \text{Vector Double} \rightarrow \text{Vector Double} \rightarrow \text{Double}$
 $\text{dotp } v \ w = \text{sum } (\text{zipWith } (*) \ v \ w)$

$\text{zipWith} :: (a \rightarrow b \rightarrow c) \rightarrow \text{Vector } a \rightarrow \text{Vector } b \rightarrow \text{Vector } c$
 $\text{zipWith } f \ v \ w = \text{unstream } (\text{zipWith}_{\text{stream}} \ f \ (\text{stream } v) \ (\text{stream } w))$

$\text{sum} :: \text{Num } a \Rightarrow \text{Vector } a \rightarrow a$
 $\text{sum } v = \text{foldl}'_{\text{stream}} \ 0 \ (+) \ (\text{stream } v)$

Fusing vector dot product

```
dotp :: Vector Double → Vector Double → Double
dotp ≡ sum (zipWith (*) v w)
      ≡ foldl' stream 0 (+) (stream (unstream
      (zipWith stream (+) (stream v) (stream w))))
      ≡ foldl' stream 0 (+)
      (zipWith stream (+) (stream v) (stream w))
```

Fusing vector dot product: $\text{foldl}'_{\text{stream}}$

$\text{foldl}'_{\text{stream}} :: (a \rightarrow b \rightarrow a) \rightarrow a \rightarrow \text{Stream } b \rightarrow a$

$\text{foldl}'_{\text{stream}} f z (\text{Stream step } s) = \text{loop } z s$

where

$\text{loop } z s = z \text{ 'seq'}$

case step s **of**

Yield x s' $\rightarrow \text{loop } (f z x) s'$

Skip s' $\rightarrow \text{loop } z s'$

Done $\rightarrow z$

Fusing vector dot product: `zipWithstream`

$\text{zipWith}_{\text{stream}} :: (a \rightarrow b \rightarrow c) \rightarrow \text{Stream } a \rightarrow \text{Stream } b \rightarrow \text{Stream } c$

$\text{zipWith}_{\text{stream}} f (\text{Stream step}_a sa na) (\text{Stream step}_b sb nb) =$
 $\text{Stream step } (sa, sb, \text{Nothing}) (\min na nb)$

where

$\text{step } (sa, sb, \text{Nothing}) =$

case $\text{step}_a sa$ **of**

$\text{Yield } x sa' \rightarrow \text{Skip } (sa', sb, \text{Just } x)$

$\text{Skip } sa' \rightarrow \text{Skip } (sa', sb, \text{Nothing})$

$\text{Done} \rightarrow \text{Done}$

$\text{step } (sa, sb, \text{Just } x) =$

case $\text{step}_b sb$ **of**

$\text{Yield } y sb' \rightarrow \text{Yield } (f x y) (sa, sb', \text{Nothing})$

$\text{Skip } sb' \rightarrow \text{Skip } (sa, sb', \text{Just } x)$

$\text{Done} \rightarrow \text{Done}$

dotp after inlining etc.

$\text{dotp} (\text{Vector } n \ u) (\text{Vector } m \ v) = \text{loop } 0.0 \ 0 \ 0$

where

$\text{loop } z \ i \ j$

$\mid i < n \wedge j < m = \text{loop } (z + u !! i * v !! j) \ (i + 1) \ (j + 1)$

$\mid \text{otherwise} \quad = z$

dotp inner loop

```
.LBB2_3:
    movsd    (%rcx), %xmm0
    mulsd    (%rdx), %xmm0
    addsd    %xmm0, %xmm1
    addq     $8, %rcx
    addq     $8, %rdx
    decq     %rax
    jne     .LBB2_3
```

Stream fusion limitations

data Stream a **where**

Stream :: (s → Step s a) → s → Int → Stream a

data Step s a = Yield a s

| Skip s

| Done

- ▶ How to support SSE instructions in this framework?
- ▶ Efficient implementations of append, replicate?

Generalizing Stream Fusion

- ▶ No single “best” stream representation.

Generalizing Stream Fusion

- ▶ No single “best” stream representation.
- ▶ *Bundle* multiple stream representations together.

```
data Bundle a = Bundle  
  { sSize    :: Size  
  , sElems   :: Stream a  
  , sChunks  :: Stream (Chunk a)  
  , sMultis  :: Multis a  
  }
```

Generalizing Stream Fusion

- ▶ No single “best” stream representation.
- ▶ *Bundle* multiple stream representations together.
- ▶ Stream consumer chooses most advantageous representation.

```
data Bundle a = Bundle
  { sSize    :: Size
  , sElems   :: Stream a
  , sChunks  :: Stream (Chunk a)
  , sMultis  :: Multis a
  }
```

Generalizing Stream Fusion

- ▶ No single “best” stream representation.
- ▶ *Bundle* multiple stream representations together.
- ▶ Stream consumer chooses most advantageous representation.
- ▶ Identical semantics, different cost models.

```
data Bundle a = Bundle
  { sSize    :: Size
  , sElems   :: Stream a
  , sChunks  :: Stream (Chunk a)
  , sMultis  :: Multis a
  }
```


Using SSE instructions in GHC

- ▶ Add support for SSE primitives to code generator(s).
- ▶ Add boxed SSE types: associated type `Multi a`.
- ▶ Fix the register allocator.
- ▶ High-level interface—via generalized streams and the vector library.

Stream representation for SSE computation

data Either a b = Left a | Right b

type Multis a = Stream (Either a (Multi a))

Stream representation for SSE computation

data $\text{Either } a \ b = \text{Left } a \mid \text{Right } b$

type $\text{Multis } a = \text{Stream } (\text{Either } a \ (\text{Multi } a))$

- Stream consumer has to take what it can get.

Stream representation for SSE computation

data Either a b = Left a | Right b

type Multis a = Stream (Either a (Multi a))

- ▶ Stream consumer has to take what it can get.
- ▶ Fine for map or fold...

Stream representation for SSE computation

data Either a b = Left a | Right b

type Multis a = Stream (Either a (Multi a))

$\text{msum}_s :: (\text{Num } a, \text{Num } (\text{Multi } a)) \Rightarrow \text{Multis } a \rightarrow a$

$\text{msum}_s (\text{Stream step } s _) = \text{loop } 0.0 \ 0.0 \ s$

where

loop summ sum1 s =

case step s **of**

Yield (Left x) s' \rightarrow loop summ (sum1 + x) s'

Yield (Right y) s' \rightarrow loop (summ + y) sum1 s'

Skip s' \rightarrow loop summ sum1 s'

Done \rightarrow multifold (+) sum1 summ

Stream representation for SSE computation

data Either a b = Left a | Right b

type Multis a = Stream (Either a (Multi a))

What about zipWith?

Stream representation for SSE computation, mark II

data Multis a **where**

Multis :: (s → Step s (Multi a))
→ (s → Step s a)
→ s
→ Multis a

- ▶ Stream consumer gets to choose what it gets.
- ▶ Works for zipWith, map, fold...
- ▶ What happens when appending two streams?

Stream representation for SSE computation

data MultiStream a **where**

MultiStream :: (s → Step s (Multi a))
 → (s → Step s a)
 → s
 → MultiStream a

- But stream operations may “mix” scalars and Multi's, e.g., append.

Stream representation for SSE computation

data MultiStream a **where**

MultiStream :: (s → Step s (Multi a))
 → (s → Step s a)
 → s
 → MultiStream a

- ▶ But stream operations may “mix” scalars and Multi's, e.g., append.
- ▶ Some operations, like fold, are agnostic to “mixing.” Others, like zipWith, are not.

A more difficult challenge

```
double ddotp(double* u, double* v, int n)
{
    union d2v d2s = {0.0, 0.0};
    double      s;
    int          i;
    int          m = n & (~VECTOR_SIZE);

    for (i = 0; i < m; i += VECTOR_SIZE)
        d2s.v += ((*((v2sd*)
                    (u+i))) * ((*((v2sd*)
                                (v+i)))));

    s = d2s.d[0] + d2s.d[1];

    for (; i < n; ++i)
        s += u[i] * v[i];

    return s;
}
```

SIMD dot product in Haskell

```
dotp :: Vector Double → Vector Double → Double  
dotp v w = msum (mzipWith (*) v w)
```

SIMD dot product in Haskell

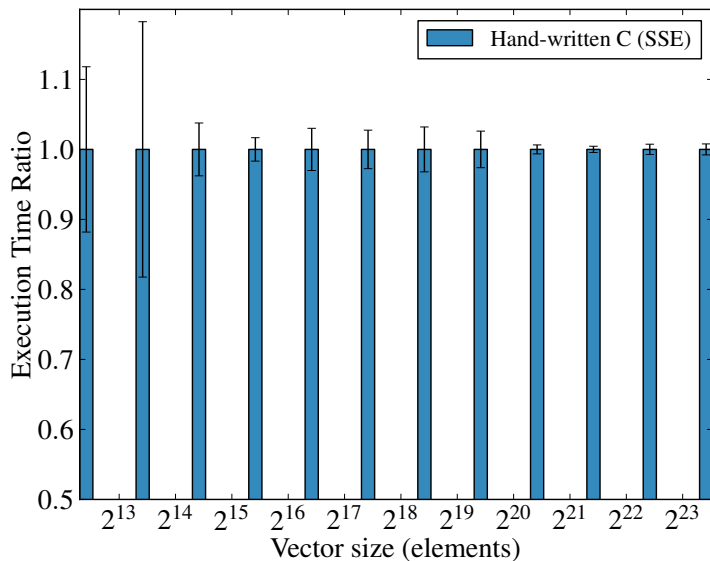
`dotp :: Vector Double → Vector Double → Double`
`dotp v w = msum (mzipWith (*) v w)`

`mfold' :: (Num a, Num (Multi a))`
`⇒ (∀b.Num b ⇒ b → b → b)`
`→ a → Vector a → a`

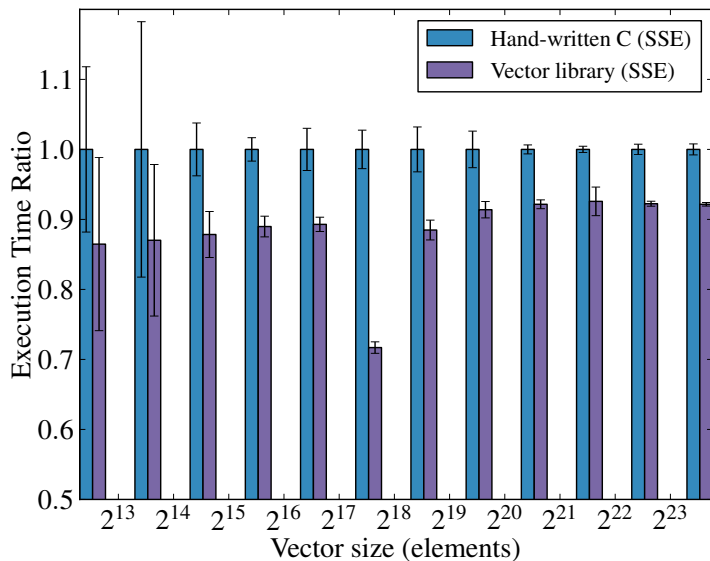
`mzipWith :: (Num a, Num (Multi a))`
`⇒ (∀b.Num b ⇒ b → b → b)`
`→ Vector a → Vector a → Vector a`

`msum :: (Num a, Num (Multi a)) ⇒ Vector a → a`
`msum = mfold' (+) 0`

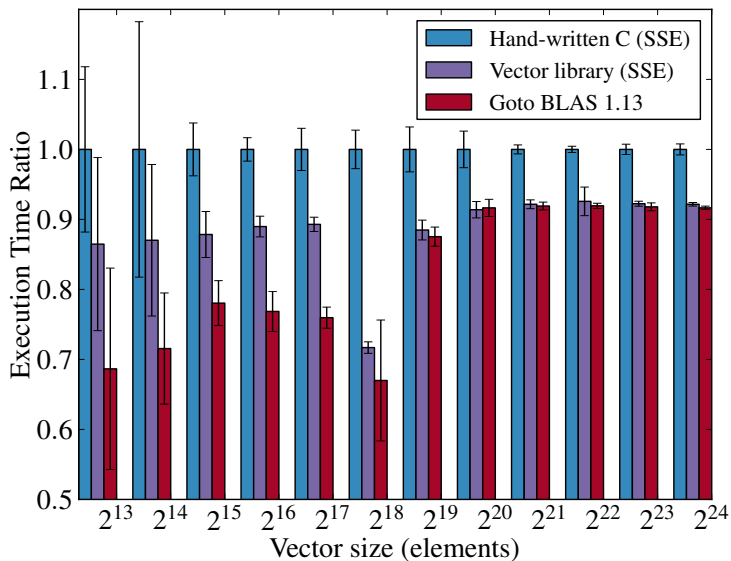
dotp performance



dotp performance: Haskell beats C



dotp performance: Haskell beats C



Haskell inner loop

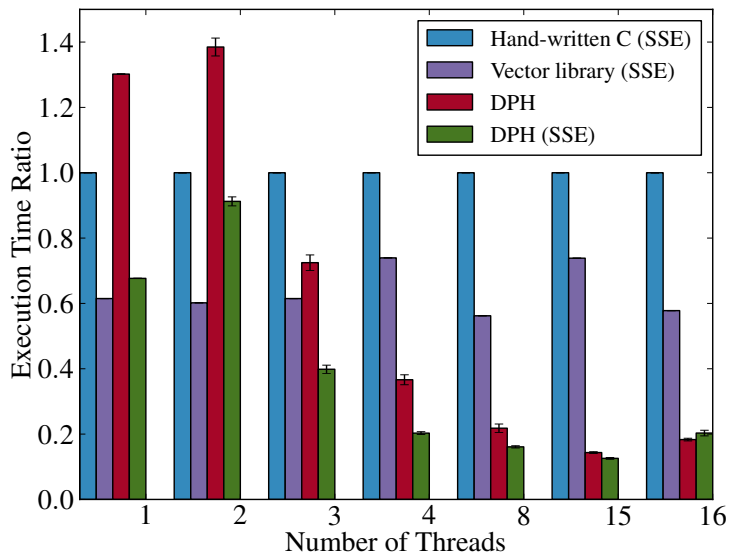
```
dotp :: Vector Double → Vector Double → Double  
dotp v w = msum (mzipWith (*) v w)
```

Haskell inner loop

`dotp :: Vector Double → Vector Double → Double`
`dotp v w = msum (mzipWith (*) v w)`

```
.LBB5_5:  
    prefetcht0 (%rdx)  
    movupd     -1408(%rdx), %xmm2  
    prefetcht0 (%rsi)  
    movupd     -1408(%rsi), %xmm1  
    mulpd      %xmm2, %xmm1  
    addpd      %xmm1, %xmm0  
    addq       $16, %rdx  
    addq       $16, %rsi  
    addq       $2, %rax  
    cmpq       %rcx, %rax  
    jl         .LBB5_5
```

Parallelization for free



Abstraction without cost: radial basis function

$$K(\vec{x}, \vec{y}) = e^{-\nu \|\vec{x} - \vec{y}\|^2}$$

Abstraction without cost: radial basis function

$$K(\vec{x}, \vec{y}) = e^{-\nu \|\vec{x} - \vec{y}\|^2}$$

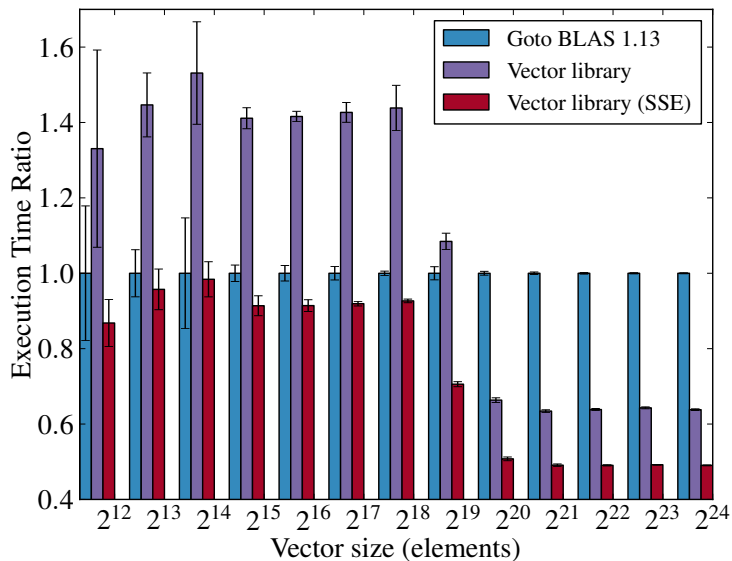
```
rbf :: Double → Vector Double → Vector Double → Double  
rbf ν v w = exp (−ν * msum (mzipWith norm v w))
```

where

```
norm x y = square (x − y)
```

```
square x = x * x
```

rbf performance



Generalized Stream Fusion Conclusions

- ▶ Key idea: simultaneously maintain multiple representations with different cost models.
- ▶ Can make use of bulk memory operations like `memcpy` and `memset` (in draft paper).
- ▶ Allows us to take advantage of SSE instructions from high-level code.
- ▶ Use of SSE instructions in DPH requires no code modifications.
- ▶ Not limited to Haskell.